

CS633 Project: Parallel Debugger

Milind Luthra (150363) Subhdeep Saha (150732)

15 March 2019

Table of Contents

- 1 Overview
 - Motivation
 - Related Work
 - Idea
- 2 Implementation
 - parallel-debugger-client (pd-client)
 - parallel-debugger-server (pd-server)
- 3 Features
 - General Debugging Features
 - MPI Specific Features
- 4 Timeline
- 5 References

Table of Contents

1 Overview

- Motivation
- Related Work
- Idea

2 Implementation

- parallel-debugger-client (pd-client)
- parallel-debugger-server (pd-server)

3 Features

- General Debugging Features
- MPI Specific Features

4 Timeline

5 References

Motivation

```
=====
- BAD TERMINATION OF ONE OF YOUR APPLICATION PROCESSES
- PID 21523 RUNNING AT hostname
- EXIT CODE: 6
- CLEANING UP REMAINING PROCESSES
- YOU CAN IGNORE THE BELOW CLEANUP MESSAGES
=====
```

```
YOUR APPLICATION TERMINATED WITH THE EXIT STRING: Aborted (signal 6)
This typically refers to a problem with your application.
Please see the FAQ page for debugging suggestions
```

Table of Contents

1 Overview

- Motivation
- Related Work
- Idea

2 Implementation

- parallel-debugger-client (pd-client)
- parallel-debugger-server (pd-server)

3 Features

- General Debugging Features
- MPI Specific Features

4 Timeline

5 References

- Debuggers already in use to debug large parallel applications.

Allinea DDT and Totalview

- Debuggers already in use to debug large parallel applications.
- Both have a rich feature set and GUIs.

Allinea DDT and Totalview

- Debuggers already in use to debug large parallel applications.
- Both have a rich feature set and GUIs.
- However, both are proprietary, commercial software.

Allinea DDT and Totalview

- Debuggers already in use to debug large parallel applications.
- Both have a rich feature set and GUIs.
- However, both are proprietary, commercial software.

Allinea DDT and Totalview

- Debuggers already in use to debug large parallel applications.
- Both have a rich feature set and GUIs.
- However, both are proprietary, commercial software.
- Restrictive licenses (locked to one node, or four processes etc) and high cost (a few hundred dollars).

Allinea DDT and Totalview

- Debuggers already in use to debug large parallel applications.
- Both have a rich feature set and GUIs.
- However, both are proprietary, commercial software.
- Restrictive licenses (locked to one node, or four processes etc) and high cost (a few hundred dollars).
- Can't be extended any further.

Using XTerm and GDB

- Possible Idea: For n processes, launching n XTerm instances with gdb.

Using XTerm and GDB

- Possible Idea: For n processes, launching n XTerm instances with gdb.
- Each terminal can be used to debug the individual processes.

Using XTerm and GDB

- Possible Idea: For n processes, launching n XTerm instances with gdb.
- Each terminal can be used to debug the individual processes.
- `mpirun -n 4 xterm -e gdb ./test`

Problems with XTerm + GDB

```
mpiexec -n 30 xterm -e gdb ./test
```

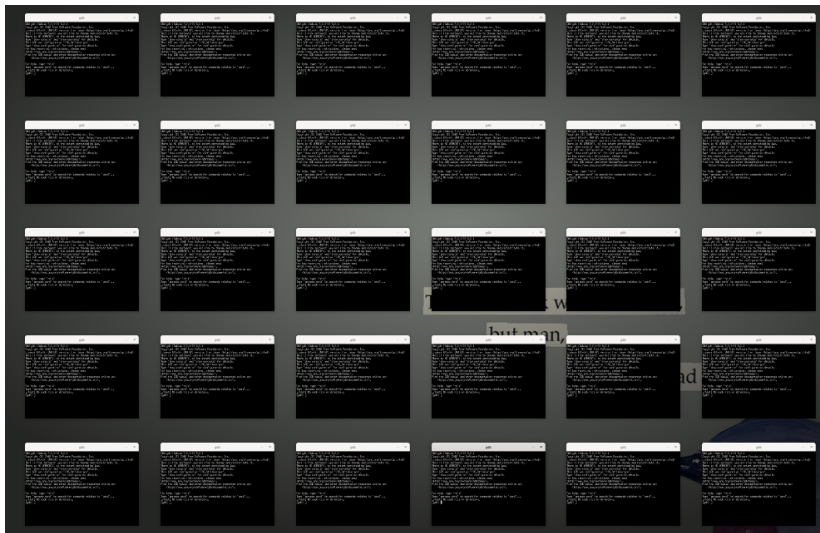


Table of Contents

1 Overview

- Motivation
- Related Work
- Idea

2 Implementation

- parallel-debugger-client (pd-client)
- parallel-debugger-server (pd-server)

3 Features

- General Debugging Features
- MPI Specific Features

4 Timeline

5 References

Basic Idea

- The basic idea is to launch a number of clients on nodes using `mpiexec`.

Basic Idea

- The basic idea is to launch a number of clients on nodes using `mpiexec`.
- Each client instance will run a `gdb` instance with the program to be debugged.

Basic Idea

- The basic idea is to launch a number of clients on nodes using `mpiexec`.
- Each client instance will run a `gdb` instance with the program to be debugged.
- All the instances of `gdb` are controlled using a single, centralized interface.

Basic Idea

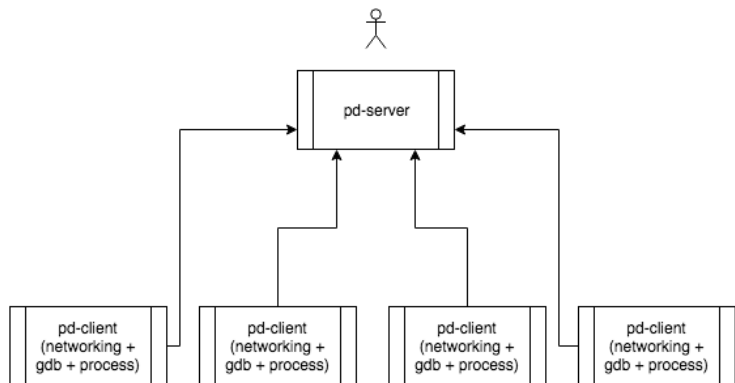


Table of Contents

- 1 Overview
 - Motivation
 - Related Work
 - Idea
- 2 Implementation
 - parallel-debugger-client (pd-client)
 - parallel-debugger-server (pd-server)
- 3 Features
 - General Debugging Features
 - MPI Specific Features
- 4 Timeline
- 5 References

Table of Contents

1 Overview

- Motivation
- Related Work
- Idea

2 Implementation

- parallel-debugger-client (pd-client)
- parallel-debugger-server (pd-server)

3 Features

- General Debugging Features
- MPI Specific Features

4 Timeline

5 References

- Handles the networking of an individual client to the server.

- Handles the networking of an individual client to the server.
 - Sends request to the server using **TCP sockets**.

- Handles the networking of an individual client to the server.
 - Sends request to the server using **TCP sockets**.
- Loads the target binary `./test` and instrument it using LD preload.

- Handles the networking of an individual client to the server.
 - Sends request to the server using **TCP sockets**.
- Loads the target binary `./test` and instrument it using LD preload.
- Acts as a frontend to the GDB interpreter `mi2`.

- Handles the networking of an individual client to the server.
 - Sends request to the server using **TCP sockets**.
- Loads the target binary `./test` and instrument it using LD preload.
- Acts as a frontend to the GDB interpreter `mi2`.
 - Conveys the instructions from the server to individual clients.

Machine Interface (mi)

- Currently GDB supports multiple command interpreters, and some of them are used to make user interface for the GDB

Machine Interface (mi)

- Currently GDB supports multiple command interpreters, and some of them are used to make user interface for the GDB
- **mi2** is a line based machine oriented text interface to GDB

Machine Interface (mi)

- Currently GDB supports multiple command interpreters, and some of them are used to make user interface for the GDB
- **mi2** is a line based machine oriented text interface to GDB
- The default interpreter is the **Console Interpreter**. For a different interpreter user can use
- `gdb .test --interpreter=(console|mi|mi2)`

mi2

- `-file-exec-and-symbols`
`./test`

console

- `file ./test`

mi2

- `-file-exec-and-symbols`
`./test`
- `-break-insert main`

console

- `file ./test`
- `break main`

mi2

- `-file-exec-and-symbols`
`./test`
- `-break-insert main`
- `-exec-run`

console

- `file ./test`
- `break main`
- `run`

mi2

- `-file-exec-and-symbols`
`./test`
- `-break-insert main`
- `-exec-run`
- `-exec-continue`

console

- `file ./test`
- `break main`
- `run`
- `continue`

Table of Contents

1 Overview

- Motivation
- Related Work
- Idea

2 Implementation

- parallel-debugger-client (pd-client)
- parallel-debugger-server (pd-server)

3 Features

- General Debugging Features
- MPI Specific Features

4 Timeline

5 References

- Ensures that all the clients are connected to it.

- Ensures that all the clients are connected to it.
- Maintains an in-memory `map[rank int] conn *net.Conn` to store the clients that are already connected.

Table of Contents

- 1 Overview
 - Motivation
 - Related Work
 - Idea
- 2 Implementation
 - parallel-debugger-client (pd-client)
 - parallel-debugger-server (pd-server)
- 3 Features
 - General Debugging Features
 - MPI Specific Features
- 4 Timeline
- 5 References

Table of Contents

1 Overview

- Motivation
- Related Work
- Idea

2 Implementation

- parallel-debugger-client (pd-client)
- parallel-debugger-server (pd-server)

3 Features

- General Debugging Features
- MPI Specific Features

4 Timeline

5 References

General Debugging Features

- The main feature is to expose all the power of gdb in a user-friendly way (through a TUI).

General Debugging Features

- The main feature is to expose all the power of gdb in a user-friendly way (through a TUI).
- Example: Set a breakpoint on `MPI_Allgather` in *each* process.

General Debugging Features

- The main feature is to expose all the power of gdb in a user-friendly way (through a TUI).
- Example: Set a breakpoint on `MPI_Allgather` in *each* process.
- (pdb) `break MPI_Allgather`

General Debugging Features

- The main feature is to expose all the power of gdb in a user-friendly way (through a TUI).
- Example: Set a breakpoint on `MPI_Allgather` in *each* process.
- `(pdb) break MPI_Allgather`
- Example: Selectively print `argc` in processes with rank 4, 5, 6 (with respect to `MPI_COMM_WORLD`).

General Debugging Features

- The main feature is to expose all the power of gdb in a user-friendly way (through a TUI).
- Example: Set a breakpoint on `MPI_Allgather` in *each* process.
- `(pdb) break MPI_Allgather`
- Example: Selectively print `argc` in processes with rank 4, 5, 6 (with respect to `MPI_COMM_WORLD`).
- `(pdb) pdb_print argc [r=4,5,6]`

General Debugging Features

- Example: List all communicators and associated processes' rank.

General Debugging Features

- Example: List all communicators and associated processes' rank.
- (pdb) `pdb_listcomm`

General Debugging Features

- Example: List all communicators and associated processes' rank.
- (pdb) `pdb_listcomm`
- Example: Selectively set breakpoint in processes in communicator 4.

General Debugging Features

- Example: List all communicators and associated processes' rank.
- (pdb) `pdb_listcomm`
- Example: Selectively set breakpoint in processes in communicator 4.
- (pdb) `pdb_break MPI_Allgather [c=4]`

Table of Contents

1 Overview

- Motivation
- Related Work
- Idea

2 Implementation

- parallel-debugger-client (pd-client)
- parallel-debugger-server (pd-server)

3 Features

- General Debugging Features
- MPI Specific Features

4 Timeline

5 References

Missing Collective Calls

- Missing a collective call in any one process of a communicator is a source of errors.

Missing Collective Calls

- Missing a collective call in any one process of a communicator is a source of errors.
- We are maintaining a communicator to processes mapping.

Missing Collective Calls

- Missing a collective call in any one process of a communicator is a source of errors.
- We are maintaining a communicator to processes mapping.
- We can use this to maintain a list of collectives pending on a process.

Missing Collective Calls

- Missing a collective call in any one process of a communicator is a source of errors.
- We are maintaining a communicator to processes mapping.
- We can use this to maintain a list of collectives pending on a process.
- Example: Listing all pending collectives.

Missing Collective Calls

- Missing a collective call in any one process of a communicator is a source of errors.
- We are maintaining a communicator to processes mapping.
- We can use this to maintain a list of collectives pending on a process.
- Example: Listing all pending collectives.
- (pdb) pdb_listcoll

Modifying Buffer During Asynchronous Calls

- A buffer passed to `MPI_Isend` cannot be written to until we're sure that it has been copied.

Modifying Buffer During Asynchronous Calls

- A buffer passed to `MPI_Isend` cannot be written to until we're sure that it has been copied.
- Usually done using `MPI_Wait` or `MPI_Test`.

Modifying Buffer During Asynchronous Calls

- A buffer passed to `MPI_Isend` cannot be written to until we're sure that it has been copied.
- Usually done using `MPI_Wait` or `MPI_Test`.
- We can track asynchronous calls using breakpoints.

Modifying Buffer During Asynchronous Calls

- A buffer passed to `MPI_Isend` cannot be written to until we're sure that it has been copied.
- Usually done using `MPI_Wait` or `MPI_Test`.
- We can track asynchronous calls using breakpoints.
- We can track buffer writes using watchpoints.

Modifying Buffer During Asynchronous Calls

- A buffer passed to `MPI_Isend` cannot be written to until we're sure that it has been copied.
- Usually done using `MPI_Wait` or `MPI_Test`.
- We can track asynchronous calls using breakpoints.
- We can track buffer writes using watchpoints.
- Hence, we can give feedback for any invalid write.

Table of Contents

- 1 Overview
 - Motivation
 - Related Work
 - Idea
- 2 Implementation
 - parallel-debugger-client (pd-client)
 - parallel-debugger-server (pd-server)
- 3 Features
 - General Debugging Features
 - MPI Specific Features
- 4 Timeline
- 5 References

Timeline

- Load and then instrument target binary in all nodes. DONE

Timeline

- Load and then instrument target binary in all nodes. DONE
- Communicate with gdb/mi. DONE

Timeline

- Load and then instrument target binary in all nodes. DONE
- Communicate with gdb/mi. DONE
- Set up client and server communication (extract rank/size from instrumented binary). DONE

Timeline

- Load and then instrument target binary in all nodes. DONE
- Communicate with gdb/mi. DONE
- Set up client and server communication (extract rank/size from instrumented binary). DONE
- Allow commands from server to all clients, and echo client console on server. DONE

- Allow sending commands from server to only specific clients TODO

Timeline

- Allow sending commands from server to only specific clients TODO
- Make a nice TUI. TODO

Timeline

- Allow sending commands from server to only specific clients TODO
- Make a nice TUI. TODO
- Track communicator data and collective calls. IF POSSIBLE

Timeline

- Allow sending commands from server to only specific clients TODO
- Make a nice TUI. TODO
- Track communicator data and collective calls. IFPOSSIBLE
- Track asynchronous call buffers. IFPOSSIBLE
- Make a nice GUI. IFANYONEVOLUNTEERS

Table of Contents

- 1 Overview
 - Motivation
 - Related Work
 - Idea
- 2 Implementation
 - parallel-debugger-client (pd-client)
 - parallel-debugger-server (pd-server)
- 3 Features
 - General Debugging Features
 - MPI Specific Features
- 4 Timeline
- 5 References

References

- [GDB/MI2 Documentation](#)
- [Go library for talking to GDB/MI](#) - Our modified version of the library.
- [Allinea DDT](#), [TotalView](#)
- [OpenMPI FAQ for Parallel Debugging](#)
- [Common MPI Programming Errors](#)